

FORGE: Force-Aware Multi-Granularity Control for Efficient Humanoid Loco-Manipulation

Anonymous Authors




Fig. 1: **FORGE** enables efficient humanoid loco-manipulation via multi-granularity control: velocity navigation, foot placement, single-/dual-arm end-effector tracking, and force-conditioned contact adaptation.

Abstract—Humanoid loco-manipulation requires coordinated locomotion, balance regulation, and manipulation under frequent contacts, making teleoperated demonstration collection difficult: a practical controller must simultaneously provide precision, stability, and contact awareness. We propose **FORGE**, a force-aware multi-granularity control system with a coarse-to-fine command space spanning velocity navigation, planar foot placement, wrist-pose tracking, and optional force-conditioning inputs. Policy training requires no full-body human motion trajectories or joint-angle references; its operator-facing interface uses sparse bilateral wrist poses relative to the pelvis, coarse motion commands, and optional desired wrist-load inputs. For sustained contact, **FORGE** combines a momentum observer with a learned correction trained from simulator-provided force labels. At deployment, it estimates low-frequency, wrist-equivalent interaction forces from proprioception and reconstructed low-level arm PD torques, without wrist force/torque sensors. A quasi-static virtual-target loop uses this estimate and an optional desired applied-force input to adapt the wrist target. A frozen locomotion prior and asymmetric residual actions stabilize manipulation-induced whole-body corrections. Experiments evaluate unified-control stability, task-space tracking, quasi-static wrist-equivalent force estimation, contact loads under matched nominal penetration commands,

and downstream diffusion-policy execution of a squat-and-lift task across payload conditions. Code and data will be released.

I. INTRODUCTION

Humanoid loco-manipulation requires coordinating manipulation with stepping, leaning, or posture adjustment under complex dynamics and frequent contacts [1]–[3]. Unlike fixed-base dual-arm robots, which benefit from stable support and mature teleoperation systems [4], humanoids must map operator intent to whole-body behavior under underactuation and changing contacts. Balance errors can compromise stability, joint errors can accumulate into end-effector drift, and rigid position control can disrupt sustained interaction. A practical controller must therefore jointly provide *Stability*, *Precision*, and *Contact Awareness*, including interaction-force estimation for regulating sustained contact.

This stable fixed-base setting has enabled scalable demonstration collection and VLA-based manipulation policies [5].

Existing approaches satisfy these requirements only in part. Full-body tracking and large-scale motor-control mod-

els [6]–[8] provide broad motion repertoires, but dense motion-reference objectives do not directly expose task-space interfaces for precise foot placement, wrist targeting, and sustained contact. Decoupled locomotion/upper-body stacks [9], [10] preserve locomotion stability but can limit task-dependent whole-body coordination. Force-aware approaches [11]–[13] demonstrate learned force estimation and force-conditioned position control, yet sensorless estimation remains difficult because external contact, friction, actuator behavior, and whole-body coupling are entangled in proprioception.

These limitations motivate **FORGE**, which *forges* a **FOR**ce-aware multi-**GR**anularity policy for **EFF**icient humanoid loco-manipulation. Its unified coarse-to-fine interface supports base velocity commands, planar foot-placement targets, Cartesian wrist-pose targets, and optional force-conditioning inputs, individually or jointly. Granularity-conditioned masking [14] activates the observations and rewards relevant to each command group. Training requires neither full-body human trajectories nor human joint-angle references; the teleoperation frontend provides only bilateral wrist $SE(3)$ poses relative to the pelvis and coarse commands, obtainable from wearable trackers, handheld VR devices, or vision. For sustained contact, a *Physics-Informed Hybrid Force Estimator* learns a simulator-supervised correction to a model-based momentum observer. At deployment, it uses arm joint states and reconstructed low-level PD torques without wrist force/torque sensors. The resulting low-frequency, wrist-equivalent interaction-force estimate is combined with an optional desired applied-force input in a quasi-static virtual-target loop. Our force-related claims therefore concern sensorless estimation and contact-load adaptation, not target-force tracking or dynamic impedance characterization.

To stabilize learning under mixed objectives, we freeze a pretrained lower-body locomotion expert and learn an asymmetric residual policy that controls the upper body while applying only corrective residual actions to the lower body. This locomotion prior preserves robust walking, reduces exploration difficulty, and permits manipulation and contact responses without requiring balance to re-emerge from end-to-end optimization.

Simulation and hardware experiments show that, relative to their ablations, masking and asymmetric residual training improve unified-control stability and task-space tracking, while the hybrid estimator reduces force-estimation error under the evaluated quasi-static contacts compared with analytical-only and learned-only estimators. A diffusion policy [15] trained on collected demonstrations further executes squat-and-lift across payload conditions.

Our contributions are a unified multi-granularity command interface trained without full-body motion references, physics-informed sensorless wrist-force estimation with contact-load adaptation, and asymmetric residual learning for stable whole-body loco-manipulation.

II. RELATED WORK

A. Whole-Body Control and Teleoperation Systems

Humanoid loco-manipulation controllers broadly follow two paradigms. Motion-reference-driven policies learn broad whole-body behaviors from retargeted or captured motions [6]–[8], [16], whereas modular systems separate learned locomotion from upper-body control [9], [10]. The former relies on dense whole-body references; the latter preserves locomotion stability but can restrict task-dependent whole-body coordination. Related work also explores masked motion generation and reusable behavior priors [14], [17], [18]. In contrast, **FORGE** trains from task-space commands without full-body human motion references and uses a frozen locomotion prior with asymmetric residual learning [19], [20]. Its teleoperation input comprises only pelvis-relative wrist poses and coarse commands, independent of whether they are obtained from wearable, handheld, or vision-based trackers.

B. Force Perception and Compliant Interaction

Humanoid controllers increasingly incorporate interaction-aware or force–position objectives [21], [22], while sensorless estimators infer wrist interaction from proprioception [11]–[13]. UniFP estimates external force from proprioceptive history and maps force–position commands to compensated position targets [11]. **FORGE** follows this force–position separation but introduces a momentum-observer prior before learned correction. It estimates low-frequency translational wrist interaction for quasi-static virtual-target adaptation, rather than full-wrench sensing or dynamic impedance control.

III. METHOD

FORGE combines a compact multi-granularity command interface with a force-aware contact pathway. Navigation, stance, and wrist-pose commands share one whole-body policy; a **Physics-Informed Hybrid Force Estimator** and virtual-target loop provide sustained-contact adaptation without wrist force/torque sensors. **Asymmetric Residual Actions** anchor this policy on a frozen locomotion prior. Fig. 2 summarizes the system.

A. Multi-Granularity Control System

One policy accepts four individually or jointly active command groups: base velocity $\mathbf{v}_{\text{cmd}}$ (**Navigation**), planar left/right footholds $\mathbf{p}_{\text{feet}} \in \mathbb{R}^{2 \times 2}$ (**Stance**), bilateral wrist poses $\mathbf{T}_{\text{wrist}} \in SE(3)^2$ (**Pose**), and optional desired robot-applied wrist forces $\mathbf{f}_{\text{cmd}} \in \mathbb{R}^{2 \times 3}$ (**Force**). The force command may be supplied by an operator or upstream model and conditions the virtual-target update in Sec. III-B; it is not a direct torque command.

During training, we sample command sets, replace observations associated with inactive groups by null values, and disable their rewards. Masking both observations and objectives prevents stale targets from conflicting with the active task. Four task-level activation bits expand into a per-group mask supplied with a six-step observation history;




Fig. 2: **Overall pipeline of FORGE.** The framework has three components: multi-granularity control over navigation, stance, pose, and optional force-conditioning inputs; a physics-informed hybrid force estimator and virtual-target contact-adaptation loop; and asymmetric residual actions anchored on a frozen locomotion prior. Policy training uses sparse task-space commands rather than full-body human motion references.

compound behaviors activate the required groups without selecting another controller or changing the 29-D action semantics. The deployed one-step signal contains 138 scalars: 3 navigation, 5 stance, 17 pose, and 6 desired-force command entries, together with 6 IMU quantities, 29 joint positions, 29 joint velocities, the previous 29-D action, and the 14-D momentum-observer residual. The explicit semantic-group mask is appended separately and is not counted in the 138 scalars.

As shown in Fig. 3, every step is partitioned into semantic groups for base motion, footholds, wrist poses, desired wrist forces, the 14-D observer residual, and proprioception. After the mask is applied, one shared MLP maps each per-step vector to a token. A two-layer Transformer (128-D latent, four heads, 256-D feed-forward layers) uses a learnable CLS token to summarize the sequence as h_{CLS} . In parallel, a 512–256–128 adaptation encoder maps the same six masked observable frames to a 64-D latent z_{adapt} ; the 512–256–128 action head consumes $[h_{CLS}, z_{adapt}]$. A 128–64 decoder maps z_{adapt} to 15 supervised outputs: a 6-D observer-force correction, bilateral wrist position (6-D), and base linear velocity (3-D). The first six outputs remain active at deployment and refine the observer estimate; wrist position and base velocity are auxiliary training targets. Simulator quantities are labels only and never actor inputs.

Stance targets are sampled within a reachable support region and expressed in the base-yaw frame. For Pose targets, we interpolate valid upper-body joint anchors, map them through wrist forward kinematics, and rotate the resulting Cartesian targets into the base-yaw frame. This procedure




Fig. 3: Actor architecture. Solid deployment paths use only the masked observable history. “Wrist Force” denotes the observer-derived residual history, not a force/torque-sensor reading or simulator force. The adaptation decoder’s 6-D correction refines the observer at deployment; the dashed privileged quantities supply training targets only. produces the reachable distributions in Fig. 4; the actor observes wrist targets rather than the sampled joints and resolves whole-body IK implicitly.

B. Physics-Informed Hybrid Force Estimator

While the Multi-Granularity Control System specifies targets, the *Force* granularity requires contact feedback. For each wrist $i \in \{L, R\}$, we distinguish the optional desired




Fig. 4: Kinematics-aware left- and right-wrist target sampling distributions.

robot-applied force $\mathbf{f}_{\text{cmd}}^i$, the wrist-equivalent environmental reaction $\mathbf{f}_{\text{env}}^i$, and its estimate $\hat{\mathbf{f}}_{\text{env}}^i$. The desired and environmental forces use the same world-frame axes but represent opposite sides of the contact pair; consequently, ideal quasi-static force balance satisfies $\mathbf{f}_{\text{env}}^i = -\mathbf{f}_{\text{cmd}}^i$. Here “wrist-equivalent” denotes the 3D translational force whose Jacobian transpose best explains the observed arm-joint disturbance under the assumed wrist-contact condition; unmodeled link contacts and whole-body coupling are not separately identified. Rather than using dedicated wrist force/torque sensors, we estimate its low-frequency, quasi-static component from arm joint states and reconstructed low-level PD torques. The estimator combines a *model-based momentum observer* [23] with a *learned correction*: the observer supplies a physically structured disturbance estimate, while the learned term compensates repeatable residual errors from the reduced arm model and actuation.

1) *Model-Based Momentum Observer*: We first estimate a joint-space disturbance residual using an arm momentum observer [23], avoiding direct acceleration differentiation. Let $\mathbf{q}^u \in \mathbb{R}^{14}$ contain the two 7-DoF arm configurations and let $\mathbf{M}^u(\mathbf{q}^u)$ assemble their arm-only mass matrices. With internal momentum state $\hat{\mathbf{p}}$ and observer residual $\mathbf{c} \in \mathbb{R}^{14}$, each control step applies

$$\hat{\mathbf{p}}_{k+1} = \hat{\mathbf{p}}_k + \Delta t(\boldsymbol{\tau}_m - \hat{\mathbf{g}} - \hat{\boldsymbol{\tau}}_f + \mathbf{c}_k), \quad (1)$$

$$\mathbf{c}_{k+1} = \beta \mathbf{c}_k + \Delta t \mathbf{K}(\mathbf{M}^u(\mathbf{q}^u) \dot{\mathbf{q}}^u - \hat{\mathbf{p}}_{k+1}). \quad (2)$$

where $\boldsymbol{\tau}_m$ is a PD-torque proxy reconstructed from commanded joint targets and measured joint states, rather than motor-current torque; $\hat{\mathbf{g}}$ and $\hat{\boldsymbol{\tau}}_f$ model gravity and smooth Coulomb friction. We set $\beta = 0.9$ and initialize $\hat{\mathbf{p}}_0 = \mathbf{c}_0 = \mathbf{0}$. The assembled arm-only CRBA matrices preserve within-arm configuration coupling but omit the floating base, waist, legs, and cross-arm coordinates. Consequently, velocity-dependent terms, excluded whole-body coupling, and actuation mismatch are not canceled analytically; they remain in the residual learned by the correction pathway. We split $\mathbf{c} = [\mathbf{c}^L; \mathbf{c}^R]$ and project each 7-D arm residual through its translational wrist Jacobian $\mathbf{J}_v^i$:

$$\hat{\mathbf{f}}_{\text{obs}}^i = (\mathbf{J}_v^i \mathbf{J}_v^{i\top} + \lambda \mathbf{I}_3)^{-1} \mathbf{J}_v^i \mathbf{c}^i, \quad i \in \{L, R\}, \quad (3)$$

where λ improves numerical stability near singularities. The result is a wrist-equivalent translational force under the assumed wrist-contact condition, not a separately identified full-body wrench. We use diagonal gain $\mathbf{K} = 20\mathbf{I}$, smooth Coulomb coefficient 1.5 N/m with $\epsilon = 0.05$ rad/s, joint-residual clipping at 20, and $\lambda = 10^{-3}$. The wrist estimate

is norm-clipped at 50 N and exponentially smoothed with coefficient 0.8; observer update and projection run at the 50 Hz policy rate.

2) *Learned Compensation with Privileged Supervision*:

The force-correction output of the adaptation decoder is denoted $\Delta \mathbf{f}_\theta(\mathbf{o}_{t-5:t})$. Rather than relearning the complete force mapping, it predicts only the discrepancy between the structured observer output and the simulator contact-force label from the same six masked frames. The final estimate and training objective are

$$\hat{\mathbf{f}}_{\text{env}} = \hat{\mathbf{f}}_{\text{obs}} + \Delta \mathbf{f}_\theta(\mathbf{o}_{t-5:t}), \quad \mathcal{L}_{\text{est}} = \mathbb{E} \left[\left\| \hat{\mathbf{f}}_{\text{env}} - \mathbf{f}_{\text{sim}} \right\|_2^2 \right]. \quad (4)$$

Here $\mathbf{f}_{\text{sim}}$ is a training-only label and never a policy input. At deployment, the decoder predicts the correction from observable history; the separate $\mathbf{f}_{\text{cmd}}$ conditions target adaptation below and is not used as the force-estimation label.

3) *Force-Conditioned Virtual-Target Adaptation*: Following the force–position separation used in UniFP [11], we use the desired applied-force input to bias a position target rather than directly commanding joint torque. For each wrist $i \in \{L, R\}$, the desired environmental reaction is $\mathbf{f}_{\text{env}}^{i,*} = -\mathbf{f}_{\text{cmd}}^i$. Around the nominal wrist target $\mathbf{p}_w^{i,*}$, we define

$$\mathbf{e}_f^i = \hat{\mathbf{f}}_{\text{env}}^i - \mathbf{f}_{\text{env}}^{i,*} = \hat{\mathbf{f}}_{\text{env}}^i + \mathbf{f}_{\text{cmd}}^i, \quad \mathbf{p}_w^{i,\text{virt}} = \mathbf{p}_w^{i,*} + \mathbf{K}_f^{-1} \mathbf{e}_f^i, \quad (5)$$

where the signed force-balance error vanishes when the desired robot-applied force is balanced by the estimated environmental reaction. The diagonal matrix $\mathbf{K}_f$ has units N/m, and $\mathbf{K}_f^{-1}$ maps the mismatch to a position offset in m/N. The two wrist updates are computed independently and then supplied jointly as bilateral pose commands to the same whole-body policy. The policy tracks $\mathbf{p}_w^{i,\text{virt}}$, so the estimated interaction changes geometric wrist targets rather than directly changing joint torques. Force estimation and target updates operate in world coordinates; policy-facing commands are rotated into the base-yaw frame. The 50 N norm bound and exponential smoothing are applied to the estimated force before its conversion into the virtual-target offset; they are not a separate hard clip on $\mathbf{p}_w^{i,\text{virt}}$. Thus, the deployment loop is closed entirely through observable arm state and commanded targets: the torque proxy drives the observer, the learned correction refines its environmental-force estimate, and the desired applied-force input together with that estimate updates the wrist target tracked by the policy. This is an algebraic, admittance-like target adaptation; because it contains no explicit inertia or damping model, it should not be interpreted as dynamic impedance control.

At each 50 Hz policy step, the observer, Jacobian projection, force-estimate clipping, smoothing, and virtual-target update execute once in this order. During training, each Cartesian component of the desired wrist-force command is sampled in $[-20, 20]$ N and smoothly ramped, held, and released, while independently generated simulator contact forces supervise the estimator correction. We use $K_f = 200$ N/m in force-command mode. Policy training uses 2,048 parallel environments, 50 simulation steps per PPO update,

and at most 20,000 updates.

C. Asymmetric Residual Whole-Body Control

Mixed locomotion, pose, and contact objectives can destabilize gait learning because unrestricted whole-body exploration can sacrifice support behavior to reduce short-horizon manipulation errors. We therefore train a lower-body expert with upper-body actions fixed to zero, freeze it, and learn a whole-body residual $\Delta \mathbf{a} = [\Delta \mathbf{a}^\ell; \Delta \mathbf{a}^u] \in \mathbb{R}^{29}$. Freezing the expert prevents gradients from manipulation and contact objectives from rewriting the nominal gait, while the trainable lower-body residual still permits command-dependent foothold and balance corrections. The residual policy receives the active multi-granularity command and learns these bounded leg corrections together with direct upper-body actions. Because the frozen action and residual are summed before one whole-body action is executed, stance adjustment and manipulation-induced posture changes remain coupled rather than being handled by independent controllers. With 15 lower-body and 14 upper-body dimensions, the executed action is

$$\mathbf{a} = [\mathbf{a}_{\text{base}}^\ell + \Delta \mathbf{a}^\ell; \Delta \mathbf{a}^u], \quad (6)$$

where $\mathbf{a}_{\text{base}}^\ell$ is the frozen locomotion action. PPO exploration uses action-standard-deviation limits of 0.2 for lower-body and 1.0 for upper-body residuals, protecting the gait prior while retaining upper-body exploration.

IV. EXPERIMENTAL EVALUATION

We evaluate four aspects of **FORGE**: (i) whether one policy supports multiple command granularities without destabilizing locomotion; (ii) how its physical-robot latency and task-space tracking compare with existing humanoid controllers; (iii) whether the hybrid estimator improves wrist-equivalent interaction-force estimation and supports contact-load adaptation; and (iv) whether the controller can serve as a downstream whole-body executor for visuomotor manipulation.

A. Comparison with Existing Humanoid Controllers

We compare **FORGE** with HOMIE [9], TWIST2 [16], and SONIC [7] on the same robot and target sequences over 20 trials per method. For TWIST2 and SONIC, we use the official implementations and checkpoints and reproduce their training pipelines locally. Command-to-action latency spans teleoperation-command emission to low-level joint-target dispatch. Foot and wrist MPE are temporal mean Euclidean target errors in planar 2D and 3D, respectively, averaged over the two limbs and then across trials. HOMIE directly maps exoskeleton commands to upper-body joints and exposes neither policy-mediated wrist targets nor explicit foothold commands, so its tracking entries are not applicable.

We additionally compare sustained contact load in simulation. Every method maintains bilateral contact with the same rigid $0.4 \times 0.4 \times 0.4$ m box and uses the same pelvis-frame targets: box center (0.25, 0, 0.25) m and left/right wrists at (0.25, ± 0.15 , 0.25) m. These targets impose matched 5 cm nominal wrist penetration. From 2 s after command onset, we

time-average the absolute normal reaction force and average both wrists. This diagnostic measures realized squeezing under a shared command; it does not isolate the estimator or virtual-target update.

B. Simulation Validation of Unified Control

Unified control, masking, and stability: Fig. 6 and Tab. V evaluate whether a single policy can coordinate velocity navigation, foot-placement control, wrist tracking, and force-conditioned contact adaptation while remaining trainable. An end-to-end 29-DoF policy trained without a locomotion prior fails to learn stable behavior, whereas the base-residual design converges. Removing residual action-std clipping causes late-stage reward collapse. The hard-switching baseline in Fig. 6 is the same ablation denoted “w/o mask” in Tab. V: inactive command observations and their objectives are not jointly gated. This ablation produces more frequent early terminations and larger velocity, foot, and wrist errors, consistent with interference between inactive command groups and active objectives. Together, the residual prior, lower-body exploration clipping, and masking mechanism support stable multi-granularity training under the evaluated command schedule.

Training protocol: All RL ablations use the same network, backbone, and environment randomization. Early training fixes yaw-rate commands at zero, smooths torque commands with EMA ($\alpha = 0.3$), and applies left-right symmetric augmentation. Task-space termination is disabled for the first 300 steps of every episode and until 2,000 global updates; afterward, we use $\delta_{\text{term}} = 0.2$ m from Fig. 7. The lower-body residual action standard deviation is clipped at 0.2.

The training reward is $R_t = \sum_i w_i r_i$, with all weights scaled by the policy period $\Delta t = 0.02$ s. Tracking terms use $\exp(-\|e\|^2/\sigma)$ unless noted otherwise, with $\sigma = 0.25$ and wrist-position scale $\sigma_w = 0.20$. Table II summarizes the task-defining rewards and principal regularizers.

Optimization, simulation, and actuation: The locomotion expert and residual stage use Adam PPO with learning rate 10^{-3} under an adaptive schedule targeting KL 0.01, discount $\gamma = 0.99$, GAE parameter $\lambda = 0.95$, clip ratio 0.2, entropy coefficient 0.01, five learning epochs, four minibatches, and gradient-norm limit 1.0. Both stages are capped at 20,000 updates. Isaac Gym advances at 200 Hz with the TGS solver (four position iterations), 0.01 m contact offset, and zero rest offset; four-step decimation yields the 50 Hz policy rate. Actions specify joint targets as $\mathbf{q}_{\text{des}} = \mathbf{q}_{\text{default}} + 0.25\mathbf{a}$. The low-level stiffness/damping pairs in N m/rad and N m/s/rad are 100/2 for hips, 150/4 for knees, 40/2 for ankles, 300/5 for the waist, and 40/1 for shoulder, elbow, and wrist joints.

Critic supervision and randomization: The actor uses six observable frames, whereas a 512–256–256 critic receives one current privileged frame containing the actor signals plus bilateral simulator wrist force (6-D), wrist position (6-D), virtual target (6-D), and base linear velocity (3-D). Auxiliary losses for wrist force, wrist position, and base velocity use relative weights 1.0/1.0/0.25, with overall adaptation-loss coefficient 1.0. Domain randomization samples added




Fig. 5: Sequential bottle placement on hardware. Four ordered frames show bottles transferred into a basket; the traces visualize vertical wrist-equivalent force responses at the marked release events and serve only as qualitative contact-event indicators.

TABLE I: Comparison with recent humanoid controllers. Hardware latency and tracking use the same robot and target sequences over 20 trials per method. Sustained load is the simulated bilateral mean absolute normal reaction under matched nominal penetration.

Method	Latency (ms)	Foot MPE (cm)	Wrist MPE (cm)	Sustained load (N)
HOMIE	27	N/A	N/A	15.48
TWIST2	197	7.87	7.76	17.16
SONIC	252	7.35	6.42	15.32
FORGE	22	5.42	4.15	8.52




Fig. 6: Unified-control ablations. (a) Asymmetric base+residual converges, unlike end-to-end 29-DoF learning. (b) Residual action-std clipping prevents late collapse. (c) Masking reduces walk→stance discontinuities.

TABLE II: Reward summary for unified-controller training. Slash-separated weights follow the listed terms; negative values denote penalties.

Group	Components (weights)	Error or objective
Command tracking	v_x/v_y (2.0/1.5), yaw/yaw integral (2.0/2.0), base height (2.0)	Base-command tracking
End-effector tasks	wrist position/quaternion (5.0/2.0), elbow/head (2.0/5.0), planar feet (5.0)	Pose and task-space tracking
Force-conditioned	virtual-wrist tracking (5.0), offset penalty (0.01), target smoothness (0.001)	$\mathbf{p}_w^{\text{virt}}$ tracking and offset regularity
Stability/posture	orientation (−3.0), waist posture (0.5), no-fly/air-time (0.75/0.3), feet/knee distance (1.0/1.5)	Balance, contact, and lateral clearance
Action/dynamics	action-rate/smoothness (−0.01/ − 0.01), power/torque (− 2×10^{-5} / − 2.5×10^{-6}), velocity/acceleration (− 10^{-4} / − 2.5×10^{-7}), foot slip (−0.25)	Motion and actuation regularization

base mass in $[-5, 10]$ kg and hand payload in $[-0.1, 0.3]$ kg, link-mass and PD-gain multipliers in $[0.9, 1.1]$, COM and body-frame offsets in $[-0.1, 0.1]$ m, friction in $[0.1, 2.5]$, and initial-joint scaling/offset in $[0.8, 1.2]/[-0.1, 0.1]$ rad. External base pushes are applied every 4 s with velocity perturbations up to 0.5 m/s in the horizontal plane and 0.2 m/s vertically.

C. Wrist-Equivalent Force Estimation and Contact-Load Adaptation

We evaluate force estimation and contact adaptation separately. Observer-only, learned-only, and hybrid variants test the estimator decomposition. A fragile-box comparison switches the desired-force-conditioning pathway on or off while keeping the robot, object, nominal wrist trajectory, and task fixed. The sustained-load column in Table I remains a




Fig. 7: Precision–stability sweep over $\delta_{\text{term}} \in \{0.1, 0.2, 0.3, 0.4\}$ m.

separate matched-command diagnostic.

Force estimation accuracy: On G1 hardware, a handheld force gauge aligned with each evaluated axis provides a digitally logged, synchronized reference but never enters the estimator (Fig. 9a). The deployed estimator uses arm joint states and reconstructed PD torques without wrist force/torque sensors. Throughout estimator evaluation, $\mathbf{f}_{\text{cmd}} = \mathbf{0}$ and every gauge load is externally applied. We repeat each load condition 10 times for staged 0–20 N loads on all axes, signed planar forces in $[-30, +30]$ N, and held-out vertical loads of 35/40/50 N. Curves in Fig. 8 are representative trials; shaded regions in Fig. 9b report one standard deviation over the ten repetitions. The learned-only baseline follows UniFP’s proprioceptive-history estimator design [11], not its complete controller.






Fig. 8: Representative single-trial real-world $x/y/z$ wrist-equivalent interaction-force estimates at staged 0, 5, 10, 15, 20 N loads. Each load is repeated 10 times with digitally synchronized force-gauge references; the gauge is used only for evaluation and is not an estimator input.



(a) Force-test setup.



(b) Planar x/y loads.

Fig. 9: Wrist-equivalent force-estimation setup and planar-load evaluation. Each applied-load condition is repeated 10 times; shaded regions denote one standard deviation, and the force-gauge stream is digitally synchronized with estimator outputs. Under external forces in $[-30, +30]$ N, the hybrid estimator records lower error than the learned-only estimator.

TABLE III: Representative staged-load trace errors in Fig. 8 (N), reported as RMSE/MAE. These time-sample errors describe the displayed trial rather than variation across repetitions.

Estimator	x axis	y axis	z axis
Hybrid	0.81/0.50	1.97/1.57	0.87/0.67
Observer-only	2.83/2.29	5.15/4.43	3.02/2.58
Learned-only	2.43/2.19	2.23/1.92	2.88/2.47

The hybrid estimate follows the gauge more closely than observer-only and learned-only variants for staged loads (Fig. 8) and records lower planar error than learned-only estimation (Fig. 9b). Table III quantifies the three estimators on the representative staged-load traces. At held-out vertical loads of 35/40/50 N, the hybrid estimator’s RMSEs are 8.28/9.66/12.35 N. These tests evaluate external-force estimation under quasi-static wrist contact, not desired-force tracking.

Contact-sensitive teleoperation: We use a controlled single-factor A/B protocol for the fragile-box lift. The position-dominant baseline disables the desired-force input and virtual-target update, whereas the force-conditioned variant enables the complete loop; policy weights, initial pose, nominal bilateral wrist trajectory, object, and lift target are unchanged. We assess two directly observable outcomes: whether the commanded lift is completed and whether grasping leaves visible foam deformation. Both variants complete the lift, while the force-conditioned run exhibits less visible deformation, linking the behavioral change to the enabled feedback pathway without claiming unmeasured contact force or material strain.

D. Hardware Tracking and Perturbation Recovery

Real-world tracking and perturbation recovery: We evaluate hardware execution across walking, squat/lift, and commanded planar foot-placement behaviors; Fig. 10 reports wrist-tracking error and pelvis acceleration during representative runs. In a separate dynamic-reaching evaluation, 3D wrist-position and planar 2D foot-position RMSEs remain below 8 cm and 10 cm, respectively. Under the evaluated 30 N wrist push and 50 cm body drag, the representative traces exhibit transient errors followed by recovery with bounded pelvis acceleration. These results characterize recovery under the specified perturbations rather than general impact-force estimation.

E. Demonstration Collection and Downstream Learning

Diffusion Policy as a validation: We use Diffusion Policy as an upstream visuomotor planner and **FORGE** as the only whole-body joint-action executor. One cross-payload model maps head-mounted RGB and task state to multi-granularity command chunks. Training pools 80 demonstrations from each of four payload conditions (320 total).

We evaluate this pipeline over 20 hardware trials for each payload on the squat-and-lift-box transport task with empty, 0.5 kg, 1 kg, and 2 kg boxes. A trial succeeds only when the robot completes the prescribed transport sequence without task termination or human intervention. As shown in Table IV, the complete pipeline succeeds in 20/20, 20/20, 18/20, and 18/20 trials, corresponding to success rates of 100%, 100%, 90%, and 90%, respectively. These results validate the upstream-planner/downstream-controller interface




Fig. 10: Pelvis acceleration and wrist-error traces across walking, squat/lift, and foot-position control under wrist-push and base-drag perturbations, with transient deviations followed by recovery.

for the evaluated task and payload range; because this table does not ablate force commands, it is not used to attribute the success-rate differences specifically to force conditioning.

TABLE V: Masking ablation metrics (mean $\pm$ standard deviation over 20 evaluation rollouts per variant); each rollout is capped at 20 s.

	w/o mask	ours
Valid rollout (s) $\uparrow$	15.69 ± 1.32	19.46 ± 0.52
Lin. velocity error (m/s) $\downarrow$	0.194 ± 0.003	0.162 ± 0.003
Planar feet MPE (m) $\downarrow$	0.182 ± 0.002	0.167 ± 0.002
Wrist MPE (m) $\downarrow$	0.115 ± 0.003	0.079 ± 0.004

V. CONCLUSION

We presented **FORGE**, a force-aware multi-granularity framework for humanoid loco-manipulation. Task-space masking, hybrid sensorless force estimation, and asymmetric residual control support coordinated whole-body execution and sustained-contact adaptation in simulation and on hardware. The present force pathway targets low-frequency, quasi-static wrist contact; it does not provide full-wrench sensing, dynamic impedance control, or force-setpoint tracking.

REFERENCES

- [1] T. He, J. Gao, W. Xiao, Y. Zhang, Z. Wang, J. Wang, Z. Luo, G. He, N. Sobanbabu, C. Pan, Z. Yi, G. Qu, K. Kitani, J. Hodgins, L. J. Fan, Y. Zhu, C. Liu, and G. Shi, "Asap: Aligning simulation and real-world physics for learning agile humanoid whole-body skills," *arXiv preprint arXiv:2502.01143*, 2025.
- [2] Q. Ben, B. Xu, K. Li, F. Jia, W. Zhang, J. Wang, J. Wang, D. Lin, and J. Pang, "Gallant: Voxel grid-based humanoid locomotion and local-navigation across 3d constrained terrains," *arXiv preprint arXiv:2511.14625*, 2025.
- [3] B. Xu, H. Weng, Q. Lu, Y. Gao, and H. Xu, "Facet: Force-adaptive control via impedance reference tracking for legged robots," 2025, arXiv:2505.06883.
- [4] Z. Fu, T. Z. Zhao, and C. Finn, "Mobile aloha: Learning bimanual mobile manipulation with low-cost whole-body teleoperation," in *Conference on Robot Learning (CoRL)*, 2024.
- [5] M. Kim, K. Pertsch, S. Karamcheti, T. Xiao, A. Balakrishna, S. Nair, R. Rafailov, E. Foster, G. Lam, P. Sanketi, Q. Vuong, T. Kollar, B. Burchfiel, R. Tedrake, D. Sadigh, S. Levine, P. Liang, and C. Finn, "Openvla: An open-source vision-language-action model," *arXiv preprint arXiv:2406.09246*, 2024.
- [6] T. He, Z. Luo, X. He, W. Xiao, C. Zhang, W. Zhang, K. Kitani, C. Liu, and G. Shi, "Omni2o: Universal and dexterous human-to-humanoid whole-body teleoperation and learning," 2024.
- [7] Z. Luo, Y. Yuan, T. Wang, C. Li, S. Chen, F. Castañeda, Z.-A. Cao, J. Li, D. Minor, Q. Ben, X. Da, R. Ding, C. Hogg, L. Song, E. Lim, E. Jeong, T. He, H. Xue, W. Xiao, Z. Wang, S. Yuen, J. Kautz, Y. Chang, U. Iqbal, L. Fan, and Y. Zhu, "Sonic: Supersizing motion

TABLE IV: Success rate over 20 hardware trials per payload of one cross-payload Diffusion Policy with **FORGE** on the squat-and-lift-box transport task. Training pools 80 demonstrations from each payload condition.

Payload	Empty	0.5 kg	1 kg	2 kg
Success rate	100%	100%	90%	90%

- tracking for natural humanoid whole-body control," *arXiv preprint arXiv:2511.07820*, 2025.
- [8] Q. Liao, T. E. Truong, X. Huang, Y. Gao, G. Tevet, K. Sreenath, and C. K. Liu, "Beyondmimic: From motion tracking to versatile humanoid control via guided diffusion," 2025, arXiv:2508.08241.
- [9] Q. Ben, F. Jia, J. Zeng, J. Dong, D. Lin, and J. Pang, "HOMIE: humanoid loco-manipulation with isomorphic exoskeleton cockpit," *CoRR*, vol. abs/2502.13013, 2025.
- [10] Y. Li, Y. Lin, J. Cui, T. Liu, W. Liang, Y. Zhu, and S. Huang, "CLONE: closed-loop whole-body humanoid teleoperation for long-horizon tasks," *CoRR*, vol. abs/2506.08931, 2025.
- [11] P. Zhi, P. Li, J. Yin, B. Jia, and S. Huang, "Learning unified force and position control for legged loco-manipulation," *arXiv preprint arXiv:2505.20829*, 2025.
- [12] Q. Lu, Y. Feng, B. Shi, M. Pisen, Z. Bao, and C. K. Liu, "Gentlehumanoid: Learning upper-body compliance for contact-rich human and object interaction," 2025, arXiv:2511.04679.
- [13] H. Wang, W. Zhang, R. Yu, T. Huang, J. Ren, F. Jia, Z. Wang, X. Niu, X. Chen, J. Chen, Q. Chen, J. Wang, and J. Pang, "Physshi: Towards a real-world generalizable and natural humanoid-scene interaction system," 2025, arXiv:2510.11072.
- [14] C. Tessler, Y. Guo, O. Nabati, G. Chechik, and X. B. Peng, "Masked-mimic: Unified physics-based character control through masked motion inpainting," *ACM Trans. Graph.*, vol. 43, no. 6, pp. 209:1–209:21, 2024.
- [15] C. Chi, S. Feng, Y. Du, Z. Xu, E. Cousineau, B. Burchfiel, and S. Song, "Diffusion policy: Visuomotor policy learning via action diffusion," in *Proceedings of Robotics: Science and Systems (RSS)*, 2023.
- [16] Y. Ze, S. Zhao, W. Wang, A. Kanazawa, R. Duan, P. Abbeel, G. Shi, J. Wu, and C. K. Liu, "Twist2: Scalable, portable, and holistic humanoid data collection system," *arXiv preprint arXiv:2511.02832*, 2025.
- [17] C. Tessler, Y. Jiang, E. Coumans, Z. Luo, G. Chechik, and X. B. Peng, "Maskedmanipulator: Versatile whole-body manipulation," 2025, arXiv:2505.19086.
- [18] Y. Li, Z. Luo, T. Zhang, C. Dai, A. Kanervisto, A. Tirinzoni, H. Weng, K. Kitani, M. Guzek, A. Touati, A. Lazaric, M. Pirodda, and G. Shi, "Bfm-zero: A promptable behavioral foundation model for humanoid control using unsupervised reinforcement learning," 2025. [Online]. Available: <https://arxiv.org/abs/2511.04131>
- [19] N. Rudin, D. Hoeller, P. Reist, and M. Hutter, "Learning to walk in minutes using massively parallel deep reinforcement learning," in *Conference on Robot Learning*. PMLR, 2022, pp. 91–100.
- [20] K. Li, P. Li, T. Liu, Y. Li, and S. Huang, "Maniptrans: Efficient dexterous bimanual manipulation transfer via residual learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025, pp. 6991–7003.
- [21] Y. Zhang, Y. Yuan, P. Gurunath, I. Gupta, S. Omidshafiei, A.-a. Aghamohammadi, M. Vazquez-Chanlatte, L. Pedersen, T. He, and G. Shi, "Falcon: Learning force-adaptive humanoid loco-manipulation," *arXiv preprint arXiv:2505.06776*, 2025.
- [22] S. Chen, Z. ang Cao, Z. Luo, F. Castañeda, C. Li, T. Wang, Y. Yuan, L. J. Fan, C. K. Liu, and Y. Zhu, "Chip: Adaptive compliance for humanoid control through hindsight perturbation," 2025, arXiv:2512.14689.
- [23] A. De Luca and R. Mattone, "Sensorless robot collision detection and hybrid force/motion control," in *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2005, pp. 999–1004.